DEBANJAN DAS

Senior Site Reliability Engineer

itsdebanjandas@gmail.com | itsme.deb1995@gmail.com | +91-7903963451

Bengaluru, India | linkedin.com/in/debanjan0212 | debanjansre.online | github.com/debanjan0212

PROFESSIONAL SUMMARY

Senior Site Reliability Engineer with 8+ years across AWS, Azure and GCP in observability, platform engineering, infrastructure as code and large-scale migrations. At Hitachi Digital Services, built the OpenTelemetry observability platform for an agentic AI product now onboarded by 5 clients, and leading its GCP build from scratch as the product goes multi-cloud across Azure, AWS and GCP. Earlier: a zero-downtime migration of 800+ servers to private EKS, a 99.8% success rate across 200+ production deployments, and a DataDog APM/RUM rollout across 200+ repositories for Toyota Motor North America. Focused on SLOs, incident management, toil reduction and secure delivery, and builds open-source agentic AI tools for SRE.

TECHNICAL SKILLS

Cloud & Containers: AWS (EC2, S3, EKS, ECS Fargate, Lambda, VPC, IAM, CloudFront, SNS/SQS, API Gateway, Route 53, ALB/NLB), Azure (AKS, Container Apps, Application Gateway, Confluent Kafka), GCP (GKE, Shared VPC, IAM), Kubernetes, Docker, Helm, OpenShift, Rancher

IaC & Automation: Terraform, Terragrunt, OpenTofu, Pulumi, CloudFormation, Ansible, Python, Bash, Linux administration

CI/CD & GitOps: Azure Pipelines, GitHub Actions, Jenkins, Argo CD, Argo Rollouts, GitLab, pre-commit (tfsec, fmt/validate)

Observability: OpenTelemetry (Collector, Contrib distro, SDK instrumentation), DataDog (APM, RUM, Logs, Monitors), Prometheus, Grafana, Loki, Tempo, Mimir (LGTM), ELK/OpenSearch, PagerDuty, ServiceNow

AI Infrastructure: Agent observability (traces, model latency), Model Context Protocol (MCP), LLM guardrails and evals, Azure OpenAI

Security: HashiCorp Vault, AWS Secrets Manager, SAST, container image scanning, vulnerability gating in CI, SOC 2/HIPAA (Vanta), RBAC

SRE Practices: SLI/SLO/SLA, error budgets, incident management (MTTD/MTTA/MTTR), blameless postmortems, fault injection, toil reduction, on-call, disaster recovery, FinOps

Data: PostgreSQL, MySQL, Aurora, MongoDB, Redis, Kafka

PROFESSIONAL EXPERIENCE

Hitachi Digital Services **Dec 2025 – Present**

Senior Site Reliability Engineer | Bengaluru, India

Operational Agentic AI (OAA) Platform – Azure, GCP, AWS

- Onboarded 5 clients onto the Operational Agentic AI (OAA) platform, an agentic AI product for IT operations triage, and drive continuous delivery, security and SRE practices across their environments.
- Leading the GCP build of the platform from scratch as OAA becomes cloud-agnostic across Azure, AWS and GCP: a private GKE cluster on a shared VPC, OpenTofu modules and an approval-gated pipeline mirroring the Azure baseline, with network ranges and IAM agreed with central IT; also assessed a client's GCP platform for integrating the product.
- Built a backend-agnostic observability platform with the OpenTelemetry Contrib Collector for 18 microservices and about 8 operational agents, plus Azure-managed Confluent Kafka, PostgreSQL, Redis and Application Gateway.
- Designed a central OTel Collector gateway exporting to DataDog and a bundled LGTM stack (Loki, Grafana, Tempo, Mimir), so customers route telemetry to their own backend with no code changes.
- Built the observability layer of the in-house agent SDK: end-to-end traces, metrics and structured logging with a standard event taxonomy, giving operators live visibility into agent execution and model latency across Kafka-connected agents.
- Delivered the LGTM stack on AKS as code with OpenTofu and an approval-gated Azure Pipeline (bootstrap, plan, approve, apply, destroy) with tfsec pre-commit checks; resolved architecture review findings including a cold-deploy race condition.
- Audited OpenTelemetry coverage across 32 platform repositories with a Python script on the GitHub API, and traced a metrics export failure (HTTP 400) to an endpoint override in the shared base library, which was fixed upstream.
- Leading a unified OpenTelemetry logging standard across core services (orchestrator, intent, context, agent-mesh) for consistent log structure and correct log levels, piloted in the orchestrator before wider rollout.

Toyota Motor North America (TMNA) – AWS

- Led DataDog APM and RUM instrumentation across 9 applications and 200+ repositories on Lambda, ECS Fargate and EKS Fargate, covering Java backends and Angular frontends.
- Built an email-to-DataDog ingestion pipeline processing 15,000–18,000 incident emails a day (1,000 per 8-minute cycle) with Grok parsing and reference-table enrichment, cutting triage lag from hours to near real time.
- Set up about 180 DataDog monitors integrated with ServiceNow and PagerDuty for automatic incident creation and escalation.
- Defined and tracked SLI/SLO targets and error budgets for critical workflows, and led blameless postmortems with development and support teams.
- Removed 5–6 hours of daily manual health-check toil with Python/Bash automation; wrote onboarding docs, CUJ dashboards and L1 templates to standardise observability.

Signeasy **Jan 2025 – Dec 2025**

Senior Site Reliability Engineer | Bengaluru, India

- Reduced Docker image size across 30+ microservices from 850 MB to 320 MB (62%) using multi-stage builds, Alpine bases and layer caching, lowering ECR storage cost and deploy time.
- Moved CI/CD to GitHub Actions with parallel jobs and caching, cutting build time from 18 to 9 minutes and CodeBuild cost by about 30%; added SAST, container image scanning and automated vulnerability gating.
- Led SOC 2 and HIPAA compliance implementation using Vanta with AWS, and managed secrets rotation with HashiCorp Vault and AWS Secrets Manager.
- Ran fault-injection testing (random container and ECS task termination) and took part in the annual disaster recovery exercise failing over from us-east-1 to us-west-2.
- Led migration of 30+ services from raw Kubernetes manifests to standard Helm charts and contributed to Terraform/Terragrunt IaC, improving release consistency and reducing configuration drift.
- Part of a 3-member team targeting 25% AWS cost reduction through Kong API Gateway migration and rightsizing.

TransUnion (Neustar acquisition) **Dec 2022 – Jan 2025**

Developer, SRE & Network Infrastructure | Bengaluru, India

- Led a zero-downtime migration of 800+ servers from on-premises to private EKS clusters for telecom clients including Charter Communications, Comcast and Oracle, and defined initial SLI/SLO metrics for the migrated platforms.
- Delivered 200+ production deployments at a 99.8% success rate, handling 50+ weekly releases for an Order Management System serving major US carriers.
- Set up monitoring for 50+ microservices and 100+ applications with ELK, Prometheus and Grafana; provided 24/7 on-call for telecom port-in/port-out operations.
- Led the OS migration from CentOS/RHEL to AlmaLinux across production; built Ansible playbooks that cut manual configuration work by 75%.

Tata Consultancy Services **Dec 2021 – Dec 2022**

System Engineer, Platform Automation | Bengaluru, India

- Led development of Kanryo, a GitOps-based SaaS platform on Kubernetes, Rancher, OpenShift, GitHub Actions and Argo CD, and mentored 15 engineers.
- Reduced Jenkins build time from 45 to 20 minutes (56%) through caching and Node.js/React pipeline cleanup.
- Built Azure AKS clusters with Terraform and extended Prometheus/Grafana monitoring with custom alerting.

Capgemini **Aug 2018 – Nov 2021**

Associate Consultant, Infrastructure & Observability | Mumbai, Bengaluru, Chennai

- Architected the ELK Stack migration from legacy Nagios for Royal Mail, reducing operational cost by 40% and incident resolution time by 60%.
- Led a 10-member support team for critical Royal Mail infrastructure at 99.5% SLA compliance; set up ITIL incident, change and problem management in ServiceNow.
- Designed a Kafka-based big data cluster for US mail processing, onboarding 8 enterprise clients in 6 months; introduced SonarQube and Nexus.

OPEN SOURCE PROJECTS

Argus – AI SRE agent over MCP | github.com/debanjan0212/argus-sre-agent

- Investigates Alertmanager alerts through MCP tools over Prometheus, Loki, Tempo and Kubernetes, finds the root cause and proposes the fix as a GitOps pull request; never writes to the cluster directly.
- Guardrails for read-only tool access, change limits, secret redaction and prompt-injection detection; CI eval suite of 5 recorded incidents, with a rule-based runbook mode when the LLM is unavailable.

IntentOps – cloud-agnostic platform agent | github.com/debanjan0212/intentops

- Turns one service intent (YAML or plain English) into capacity, SLO burn-rate alerts, policy checks, a per-cloud cost comparison and a canary that rolls back on SLO breach, for EKS, GKE, AKS, bare metal or kind.
- Ships as a CLI, a GitHub Action and an MCP server; adding a new cloud is one YAML profile.

agent-harness-sdk – Python SDK for operational AI agents | github.com/debanjan0212/agent-harness-sdk

- Read, analyse, decide, act loop with policy in code, human approval before any change, secret redaction, prompt-injection checks and budgets; emits OpenTelemetry traces, metrics, logs and events automatically.
- DeepSeek by default or any OpenAI-compatible model; adapters for LangChain, LangGraph, CrewAI, LlamaIndex, OpenAI Agents SDK, MCP and Harness.io; ships Incident Commander, a reference incident-management bot with a Helm chart.

EDUCATION

Vellore Institute of Technology (VIT), Vellore **2014 – 2018**

B.Tech, Electrical and Electronics Engineering | CGPA 8.48/10, First Class with Distinction

- Published a research paper on emotion recognition using machine learning and neuroscience; internships at Schneider Electric, TE Connectivity and BEML.

Languages: English (C2, proficient), Hindi (B2), Bengali (native)